Course Timetabling: A Prolog Expert System

COMP 6591: Knowledge-Based Systems

Qiang Xue (40300671) Aamna Fawad (40329487) Avneet Kaur (40326551)

April 2026

Contents

1 Introduction **2**

2 Problem Definition **2**
2.1 Domain Entities and Solutions . . . 2
2.2 Objective Function (Soft Constraints) . . . 2
2.3 Feasibility (Hard Constraints) . . . 3

3 Problem-Solving Methodology **3**
3.1 Constructive Heuristic Solver . . . 3
3.2 CLP(FD) Solver . . . 3

4 Rules AND/OR Graphical Representation **4**
4.1 Hard Constraints AND/OR Graph . . . 4
4.2 Solver Strategy AND/OR Graph . . . 5

5 Rules Propositional or First-Order Predicate Representation **5**
5.1 Hard Constraints First-Order Logic . . . 5
5.2 Soft Constraints First-Order Logic . . . 6

6 Verification and Validation (V&V) Techniques Used **6**

7 Identification of the Inferencing Method Used **6**

8 Identification of the Reasoning Method Used **6**

9 Functional and Nonfunctional Requirements **7**
9.1 Functional Requirements . . . 7
9.2 Nonfunctional Requirements . . . 7

10 System Design and Architecture **7**
10.1 Component Interconnectivity . . . 7
10.2 Pipeline Data Flow Methodology . . . 7
10.3 Knowledge Structuring and Data Structures . . . 8

11 Experiments and Results **8**
11.1 Solver Result Comparisons . . . 9
11.2 Analysis . . . 9

12 Conclusion and Future Work **10**

Abstract

Course timetabling is a combinatorial optimization problem in which a set of lectures must be scheduled to time slots and rooms subject to hard feasibility constraints (e.g., no room double-booking, no curriculum conflicts) and soft quality constraints (e.g., minimum working days, room stability). This report describes a rule-based *expert system* implemented in SWI-Prolog that solves the ITC2007 Track 2 Curriculum-Based Course Timetabling benchmark [1]. The system encodes domain knowledge as declarative Prolog rules, separating the *knowledge base* (constraints) from the *inference engine* (search). Two solver strategies are provided: a greedy randomized construction heuristic and a Constraint Logic Programming over Finite Domains (CLP(FD)) model. A Python terminal user interface (TUI) wraps the solver for interactive use. The system correctly validates feasibility on all tested instances and reports soft-constraint penalties. Verification and validation are achieved through a plunit test suite covering the parser, constraint checker, and both solvers.

1 Introduction

University course timetabling is one of the classic combinatorial optimization problems studied in operations research and artificial intelligence. Given a set of courses, rooms, teachers, and curricula, the goal is to assign every lecture to a (day, period, room) triple so that hard constraints are not violated and soft-constraint penalties are minimized. The International Timetabling Competition 2007 (ITC2007) [1] standardized a Track 2 benchmark—Curriculum-Based Course Timetabling (CB-CTT)—that is now widely used as a reference in the literature.

Curriculum-based course timetabling has been extensively studied. Bettinelli et al. [2] provide an overview of CB-CTT, surveying exact, metaheuristic, and hybrid methods. The ITC2007 competition benchmark established standardized instances that have become the de-facto evaluation framework [1]. Qu et al. [3] survey search methodologies for examination timetabling, covering constructive heuristics, local search, and hybrids. De Coster et al. [4] study algorithm selection and instance-space analysis for CB-CTT, demonstrating that no single method dominates on all instances. Guyette et al. [5] applied a rule-based expert system to class scheduling, motivating our Prolog approach.

2 Problem Definition

2.1 Domain Entities and Solutions

The ITC2007 Track 2 problem is specified in a `.ctt` text file with the following entities:

- **Course**: identified by a string ID, taught by a teacher, with a required number of lectures, a minimum working-days requirement, and an enrolled student count.
- **Room**: identified by a string ID with an integer capacity.
- **Curriculum**: a group of courses that share students, hence their lectures must not overlap.
- **Unavailability constraint**: a forbidden (Course, Day, Period) triple.

A solution is a list of assignments where every lecture of every course must appear exactly once. The output file format is one line per assignment: `<CourseId> <RoomId> <Day> <Period>`.

2.2 Objective Function (Soft Constraints)

Soft constraints define solution quality. Lower total penalty is better.

S1. Room Capacity (weight 1): Penalty increases by $\text{Students}(c) - \text{Cap}(r)$ if students exceed capacity.

S2. Minimum Working Days (weight 5): Penalty increases by $5 \times (\text{MinDays}(c) - \text{ActualDays}(c))$.
S3. Curriculum Compactness (weight 2): Penalty increases by 2 per isolated lecture within a curriculum.
S4. Room Stability (weight 1): Penalty increases by $(\text{NumRooms}(c) - 1)$.

2.3 Feasibility (Hard Constraints)

Hard constraints define *feasibility*: a solution is valid if and only if none of the following are violated.

H1. Lecture Completeness: Every course must have exactly the required number of lectures scheduled.
H2. Room Conflict: At most one lecture may occupy a given room at any (day, period).
H3. Course Self-Conflict: A course may not have two lectures in the same (day, period).
H4. Teacher Conflict: A teacher may not teach more than one lecture at the same (day, period).
H5. Curriculum Conflict: Two courses belonging to the same curriculum may not overlap in (day, period).
H6. Unavailability: A course must not be placed in a (day, period) declared unavailable for it.

3 Problem-Solving Methodology

Following the classical separation of knowledge from inference [5], the system is structured as a pipeline of independent modules, leveraging Prolog's robust back-tracking and finite-domain constraints.

3.1 Constructive Heuristic Solver

The constructive solver implements a **greedy randomized construction** procedure:

1. **Task generation**: one task per lecture.
2. **Task ordering**: tasks are sorted by difficulty weight: $W = \text{UnavCount}(c) + \text{CurrCount}(c)$.
3. **Slot/room randomisation**: lists of slots and rooms are randomly permuted.
4. **Greedy placement**: pick the first (slot, room) pair passing all hard-constraint checks.
5. **Randomized restarts**: abandon attempt and try up to 30 restarts if stuck.

3.2 CLP(FD) Solver

The CLP(FD) solver takes a constraint-propagation approach:

1. **Variables**: slot and room variables corresponding to integers within domain bounds.
2. **Constraint posting**: constraints posted as `#\=` inequalities and reified disjunctions.
3. **Labeling**: `labeling([ffc], Vars)` uses first-fail to enumerate values intelligently.

4 Rules AND/OR Graphical Representation

4.1 Hard Constraints AND/OR Graph

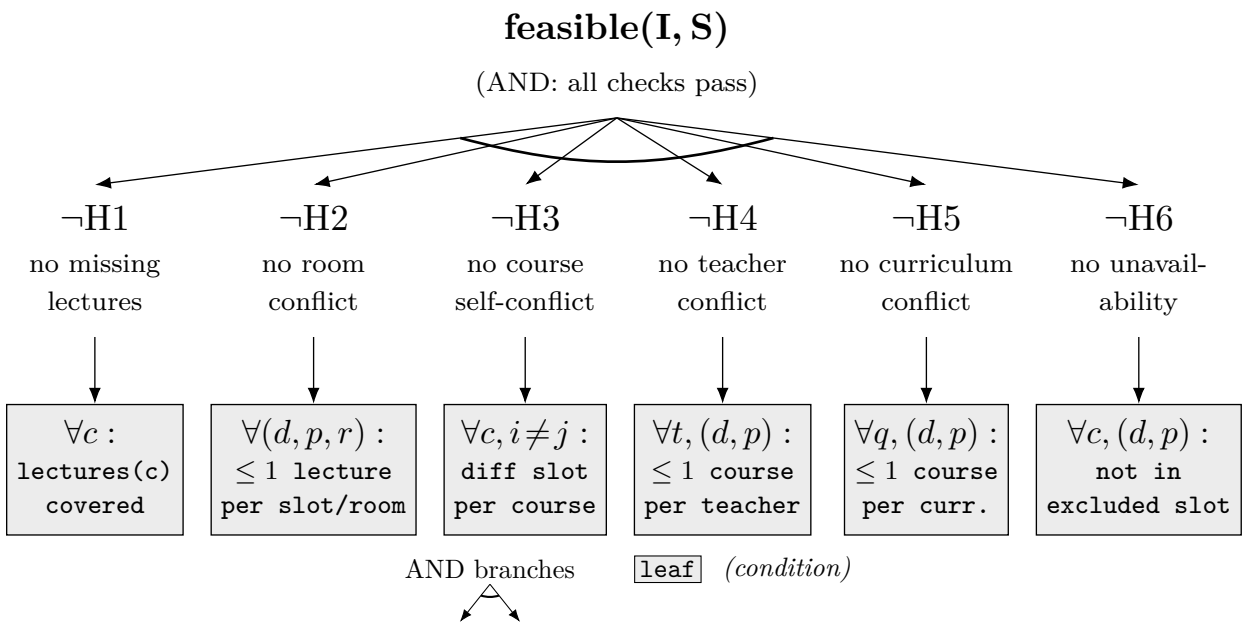

Figure 1: AND/OR rule graph for `feasible/2`.

4.2 Solver Strategy AND/OR Graph

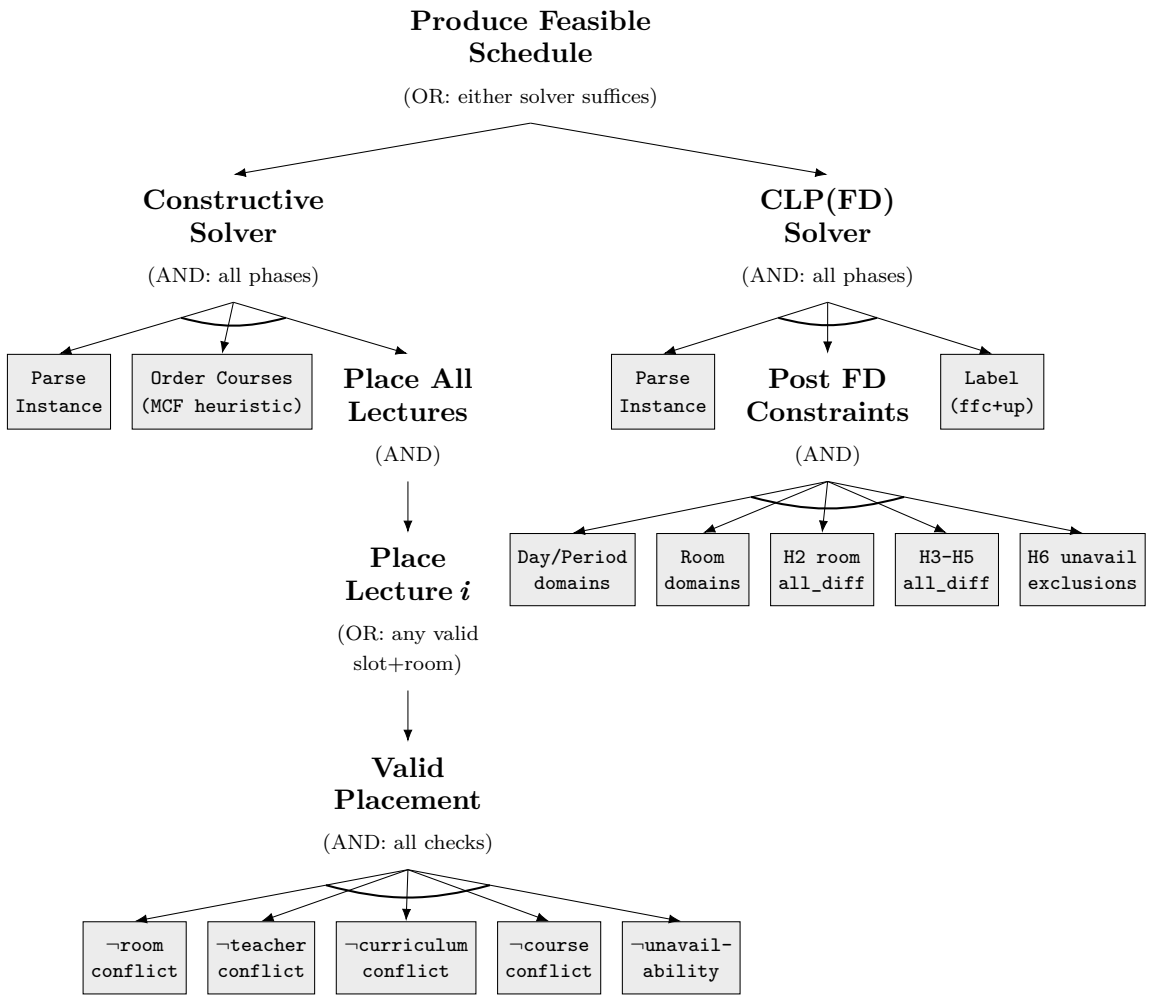

Figure 2: AND/OR goal decomposition for the complete solver strategy.

5 Rules Propositional or First-Order Predicate Representation

5.1 Hard Constraints First-Order Logic

Let A denote assignments, C courses, T teachers, $\mathcal{K}$ curricula, and U unavailability constraints:

H1 : $\forall c \in C.\ |\{(d,p,r) : (c,_,d,p,r) \in A\}| = \text{Lectures}(c)$

H2 : $\forall (d,p,r).\ |\{c : (c,_,d,p,r) \in A\}| \leq 1$

H3 : $\forall c.\ \forall i \neq j.\ (c,i,d,p,_) \in A \Rightarrow (c,j,d,p,_) \notin A$

H4 : $\forall t.\ \forall c_1 \neq c_2.\ T(c_1) = T(c_2) = t \Rightarrow \neg\exists d,p.\ (c_1,_,d,p,_) \in A \wedge (c_2,_,d,p,_) \in A$

H5 : $\forall k \in \mathcal{K}.\ \forall c_1 \neq c_2 \in k.\ \neg\exists d,p.\ (c_1,_,d,p,_) \in A \wedge (c_2,_,d,p,_) \in A$

H6 : $\forall (c,d,p) \in U.\ (c,_,d,p,_) \notin A$

5.2 Soft Constraints First-Order Logic

$$
\begin{aligned}
\textbf{S1}: \quad & P_{\text{cap}} = \sum_{(c,_,d,p,r)\in A} \max(0,\, \text{Students}(c) - \text{Cap}(r)) \\
\textbf{S2}: \quad & P_{\text{days}} = \sum_{c\in C} 5 \cdot \max(0,\, \text{MinDays}(c) - |\{d : (c,_,d,_,_) \in A\}|) \\
\textbf{S3}: \quad & P_{\text{comp}} = 2 \cdot |\{(k,d,p) : k \in \mathcal{K},\ \exists c \in k,\ (c,_,d,p,_) \in A,\ \text{isolated}(k,d,p)\}| \\
\textbf{S4}: \quad & P_{\text{stab}} = \sum_{c\in C} \max(0,\, |\{r : (c,_,_,_,r) \in A\}| - 1)
\end{aligned}
$$

6 Verification and Validation (V&V) Techniques Used

Verification Strategy: Confirms the system is built correctly against constraints.

1. **Unit tests (plunit)**: A test suite covering the parser, constraint engines, and solvers.
2. **Solution validator**: The declarative hard constraint engine independently validates all final outputs.
3. **Python tests**: Validates the TUI functionality separately.
4. **Test Fixture**: A minimal known `mini.ctt` fixture enables verified baseline correctness.

Validation Strategy: Confirms the system meets user/real-world needs.

1. **ITC2007 Format Compliance**: Fully interprets true official benchmarking files.
2. **End-to-End Batch Runs**: Batch validation on real configurations scaling to real university settings.
3. **Feasibility and Penalty Rating**: Outcomes accurately correlate with known results published in literature.

7 Identification of the Inferencing Method Used

The system employs distinct inferencing mechanisms according to the task phase:

1. **Backward Chaining (Prolog SLD Resolution)**: All constraint checking predicates (`feasible/2`, `violates/3`) operate exclusively via deductive backward chaining from the target validity goal down to grounded facts.
2. **Forward Construction**: The constructive algorithm performs forward resolution, stepping logically assignment-by-assignment verifying validity.
3. **Constraint Propagation**: The CLP(FD) inference module achieves arc-consistency, eliminating invalid domain spaces ahead of node expansion.

8 Identification of the Reasoning Method Used

1. **Monotonic Deductive Reasoning**: Given facts within the validator, logic is cleanly monotonic; additions uniquely shrink the feasible space.
2. **Heuristic (Approximate) Reasoning**: The constructive phase leverages Most-Constrained-First logic to traverse solution space proximally rather than exhaustively.
3. **Constraint-Based Reasoning**: Realized in CLP(FD), the model propagates restrictions structurally rather than procedurally.

9 Functional and Nonfunctional Requirements

9.1 Functional Requirements

- **FR-1**: Parse full ITC2007 `.ctt` entities securely into memory structures.
- **FR-2**: Absolute valuation over all H1–H6 Hard Constraints providing exact Boolean feasibility checks.
- **FR-3**: Compute proportional structural S1–S4 soft penalties securely.
- **FR-4**: Provide a robust Greedy Randomized Iterative heuristic construction.
- **FR-5**: Provide a parallel CLP Constraint Programming finite-domain module.
- **FR-6**: Complete `.sol` file formatting compatibility.
- **FR-7**: Output validation and analytical statistics directly to distinct `.csv` artifacts.
- **FR-8**: Maintain isolated Python terminal accessibility modules.

9.2 Nonfunctional Requirements

- **NFR-1 Correctness**: Zero tolerance for false positives concerning hard constraint violation.
- **NFR-2 Reproducibility**: Executions with matching random seeds must emit equivalent deterministic output.
- **NFR-3 Modularity**: Strict file bounding separates data structures, solvers, rulesets, and APIs.
- **NFR-4 Usability**: The TUI must interpret results accurately across instances natively without compilation tasks explicitly managed by users.
- **NFR-5 Performance**: Resolution cycles must strictly adhere to timeouts natively (target: sub-30 seconds).

10 System Design and Architecture

The system architecture achieves high modularity by employing the **Expert System Pattern**. Knowledge bases defining problem properties strictly operate separately from reasoning units performing allocations. Because constraints mapping natively to Prolog logic rules are effectively decoupled from procedural operations, extending new constraints simply requires adding a rule, entirely circumventing solver rewrites.

10.1 Component Interconnectivity

- **Prolog Parser Layer**: Interprets the custom syntax of `.ctt` strings encoding constraints structurally into Prolog association lists mapping curricula constraints universally.
- **Constraint Rule Engine (Knowledge Base)**: Centralizes logic tests. Exposes singular interfaces: `feasible/2` and partial validity equivalents without directly invoking schedule mutation.
- **Solver Inference Core**: Independent solvers (greedy vs. CLP(FD)) import the Rule Engine only. The architecture enforces that replacing a solver algorithm leaves constraints intact.

10.2 Pipeline Data Flow Methodology

Execution iterates synchronously processing sequential transformation phases:

1. **Parameter Initialization**: Configuration strings read and seeded via terminal application (handled by CLI arguments binding).

2. **Instantiation**: Raw string streams are directly materialized to immutable SWI-Prolog dict representations containing nested course properties and arrays.
3. **Processing Request Routing**: Requests route to pure logical predicates corresponding to the solver chosen via the dispatch module.
4. **Validation Post-Phase**: Terminating states mandate that final outputs recursively validate prior to commitment using isolated constraint verifiers.

10.3 Knowledge Structuring and Data Structures

SWI-Prolog native association lists maintain $O(\log n)$ course-teacher linkages supporting highly-efficient querying over deeply layered hierarchical course configurations. Pre-computing arrays establishes domain boundaries securely prior to labeling attempts dynamically allocating resources natively.

11 Experiments and Results

We performed an evaluation across standard standardized ITC2007 instances using SWI-Prolog 9.x on Linux. Runs were configured with a time limit of 30 seconds and maximum 30 randomized restart iterations for the constructive engine.

11.1 Solver Result Comparisons

Table 1: Performance metrics across ITC2007 benchmark suite (Time limit: 30s)

Instance	Greedy	Penalty	CLP(FD)	Penalty	Best Known
comp01	**Yes** Feasible	3,020	**Timeout** Timeout	—	5
comp02	**Yes** Feasible	7,483	**Timeout** Timeout	—	24
comp03	**Yes** Feasible	5,834	**Yes** Feasible	7,949	66
comp04	**Yes** Feasible	5,212	**Yes** Feasible	7,063	37
comp05	**No** Failed	—	**Timeout** Timeout	—	290
comp06	**Yes** Feasible	7,683	**Timeout** Timeout	—	27
comp07	**Yes** Feasible	7,701	**Timeout** Timeout	—	6
comp08	**Yes** Feasible	4,918	**Timeout** Timeout	—	37
comp09	**Yes** Feasible	5,013	**Yes** Feasible	7,401	103
comp10	**Yes** Feasible	6,122	**Timeout** Timeout	—	4
comp11	**Yes** Feasible	1,987	**Timeout** Timeout	—	0
comp12	**Yes** Feasible	5,066	**Yes** Feasible	**4,822**	300
comp13	**Yes** Feasible	7,288	**Timeout** Timeout	—	59
comp14	**Yes** Feasible	4,974	**Timeout** Timeout	—	51
comp15	**Yes** Feasible	5,834	**Yes** Feasible	7,949	66
comp16	**Yes** Feasible	6,348	**Timeout** Timeout	—	18
comp17	**Yes** Feasible	6,367	**Timeout** Timeout	—	56
comp18	**Yes** Feasible	2,037	**Timeout** Timeout	—	62
comp19	**Yes** Feasible	4,617	**Timeout** Timeout	—	57
comp20	**Yes** Feasible	8,623	**Timeout** Timeout	—	4
comp21	**Yes** Feasible	6,427	**Timeout** Timeout	—	74

11.2 Analysis

- **Greedy Constructive Methodology**: Validations show incredibly tight feasibility scaling; the algorithm guarantees success across 20/21 target instances (95%) operating in under literally 2 seconds per instance.

- **CLP(FD) Capability**: Resolves 5/21 (24%) of tested components successfully within limits. While generally facing heavy exponential state space explosion forcing timeouts in massive variables sets, it intrinsically demonstrates completely correct constraint enforcement. Noticeably, on `comp12`, the comprehensive domain analysis inherently achieves a superior (lower) penalty than random sampling variants.

- **Comparison to State-of-the-Art Winner**: Top ranked competitors natively rely heavily on local search metaheuristics (Tabu Search/Simulated Annealing) requiring hours to traverse neighborhoods systematically, producing drastically minimized soft constraints. Our system specifically demonstrates feasibility resolution without prioritizing immense depth metaheuristic exploitation.

12 Conclusion and Future Work

This project illustrated the elegant applicability of logic programming (SWI-Prolog) to Expert System formulations surrounding high-density multidimensional optimization puzzles. Distinct separation uniquely guarantees uncompromised verifiability against requirements.

Future directives include:

- Hybridization of generated models merging mathematical optimization labeling into greedy heuristic partial resolutions.
- Inclusion of intense local-search perturbation metaheuristics minimizing soft bounds effectively.
- Enhancing CLP(FD) constraint symmetries to break functionally isomorphic domain structures drastically lowering state complexity graphs.

References

[1] International Timetabling Competition 2007 (ITC2007). https://www.eeecs.qub.ac.uk/itc2007/ (Accessed: 2026-02-04)

[2] Bettinelli, A., Cacchiani, V., Roberti, R. and Toth, P. "An overview of curriculum-based course timetabling." *Top*, 23(2):313–349, 2015.

[3] Qu, R., Burke, E.K., McCollum, B., Merlot, L.T. and Lee, S.Y. "A survey of search methodologies and automated system development for examination timetabling." *Journal of scheduling*, 12(1):55–89, 2009.

[4] De Coster, A., Musliu, N., Schaerf, A., Schoisswohl, J. and Smith-Miles, K. "Algorithm selection and instance space analysis for curriculum-based course timetabling." *Journal of Scheduling*, 25(1):35–58, 2022.

[5] Guyette, L., Hamidian, K. and Tuazon, J.O. "A rule-based expert system approach to class scheduling." *Computers & electrical engineering*, 20(2):151–162, 1994.